

Parameter Based Style Transfer on Images

Universität Heidelberg

Fakultät für Physik

**Parameter Based Style Transfer on
Images**

Masters Thesis

Arthur Willy Heimbrecht

June 2nd 2020

supervised by Professors Björn Ommer and Tilman Plehn

Abstract

This is where the abstract is going to be.

Contents

Abstract	iv
Contents	v
1 Introduction	1
THEORETICAL BACKGROUND & RELATED WORK	3
2 Machine Learning and Computer Vision	5
3 Optimization and Gradient Descent	7
4 Generative Models	9
5 Style Transfer & Painterly Rendering	11
5.1 Style Transfer	11
5.2 Painterly Rendering	11
CONTRIBUTION AND EXPERIMENTS	13
6 Approach	15
6.1 Motivation	15
6.2 Neural Renderer	17
Data Set	17
Architecture	22
Training	23
Results	23
6.3 Stroke Approximation	23
Data Set	23
Weapon of Choice	24
Optimization Based Approach	28
Image Placing and Blending	32
Style Transfer	39
Optimization Details	39
Results	40
7 Ablation Experiments	41
7.1 Neural Renderer	41
7.2 Optimization Procedure	41
CONCLUSION	43
8 Conclusion	45
9 Outlook	47

APPENDIX	49
A Appendix	51
Bibliography	53

1

Introduction

Since the arrival of electric-based computers in the 20th century, the range of application for computers and algorithms has steadily grown. From the first simple arithmetic calculations over solving complex equations and simulating realistic systems to media production and funny apps for phones; computers seem to seize every aspect of everyone's life. This includes areas that stood their ground for long time against the automation that has become common for so many areas like production, communication etc.. Recent advances in Machine Learning and the ever increasing efficiency, availability and sheer power of computers has put the thought of professions like lawyers and physicians being 'future-proof' into question.

[cite definition of style](#)

This includes even the field of arts. As art comes can take pretty much every form there is and as it depends on the viewer, it is valid to ask whether art could ever be understood by anything but a human being. Nonetheless, people have tried repeatedly and proven that Machines CAN grasp broad concepts in art, even the delicate art of drawing

Evolution of Art and Style Transfer Since the late 1980's the field of computer vision has looked not only at photographs and videos but also started working on images of artworks. From early computer based analysis and identification of images that is used even today for identifying counterfeits, the field has gone as far as creating art on its own. Ever since Leon Gatys introduced his paper on 'Neural Style Transfer' in 2016, the quality of computer generated paintings has come closer and closer to being indistinguishable to an image of an original artwork. This resulted in many interpretations of given images in the style of well-known artists like van Gogh, Matisse or Picasso.

Representation Gap Yet, many of these accomplishments lack the final step towards creating authentic artworks, which is the transitions from the image domain to the actual artwork domain. As images are mere projections of the artworks they try to represent, they lack some of the original artwork's content and message. Most people would agree that it is an entirely different experience to view a painting in real life than viewing an image of this painting on a screen or sheet of paper.

For once, paint on canvas has depth to it, that can hardly be visualized by a 2D plane. Every viewing angle of an image gives away different details in the works of the artist. How much paint did the artist use? How many layers of color were applied? Which brush or technique did were used? All of this plays into an artwork's impression as much as the distribution of color does. This is the reason why people are interested in actual painted replicas of their favorite artworks, instead of just high resolution prints.

As this stands, it becomes clear why an image will always loose some of this information about the artwork. Sure, one could always make

hundreds of very high resolution images or even scan the image in 3D space, yet without advanced technology like VR headsets or the like, it still does not quite replace the real deal.

From Pixels to Strokes So why even bother then? Even with all the shortcomings which were described here, images still catch most of the essence of an artwork – at least enough to train neural networks on this. But is it possible to maybe depict artworks as what they actually are, a composition of color particles on canvas? It just so happens that there are works such as [[adobe stroke paper](#)] that introduced the idea of synthesizing brush strokes through fluid dynamics and neural networks. Unfortunately this comes with a high computational burden. So maybe one could start by at least taking a step towards this sort of representation by depicting an artwork through a set of brush strokes. Ideally, this could help to better capture the essence of an artwork if this comes closer to its real world representation and later even open the door for style transfer to be performed on this representation as well.

Structure of this work In this work an attempt at retrieving such representations for artworks is made. First, the theoretical background for such an attempt along with related work will be presented. Then the approach itself will be inferred in two parts that describe the training of a differentiable renderer and the subsequent retrieval of brush strokes from an image. This is followed by experiments as well as an ablation study of the approach. Lastly, there will be a discussion about the results and the future direction of research.

Allocation of individual contributions

THEORETICAL BACKGROUND & RELATED WORK

**Machine Learning and Computer
Vision**

2

Optimization and Gradient Descent

3

Generative Models

4

Style Transfer & Painterly Rendering

5

5.1 Style Transfer

5.1 Style Transfer 11
5.2 Painterly Rendering 11

5.2 Painterly Rendering

CONTRIBUTION AND EXPERIMENTS

6

Approach

6.1 Motivation

The basic approach of this work can be split into two steps:

1. A differentiable renderer which can generate images of brush strokes from a parameter representation.
2. An optimization procedure that iteratively approximates an image through brush strokes representations.

The two step approach can be motivated by comparing the optimization procedure to the actual process of painting an image. An artist will most likely not pick single color particles to then place them on the canvas. Instead, an artist uses a brush or other utilities (see Pollock or others) to place more paint with a single action.

This of course limits the control over each individual drop of paint but maintains enough control to still create very delicate details in paintings. This trade-off varies for different brush sizes and for this reason an artist must choose the brush size depending on the content.

An example would be the painting of a uniformly colored sky. By using a large brush size the artist can cover a lot of canvas in relatively little time as well as keep the color well distributed over the canvas because the brush will spread color more or less evenly over its footprint. On the other hand, if one were to draw a sky with the smallest brush available, not only would it take forever to paint, it would also be hard to keep the paint evenly distributed over multiple strokes.

Now, translating this onto the given problem of recreating/approximating an image through brush strokes, it would mean to limit the process to only use what we would describe as brush strokes.

An equivalent example would be the game of Tangram.

Tangram is a Chinese puzzle game that has the objective of replicating a given silhouette only with a set of 7 unique shapes. The shapes may not overlap or be cut or anything. Quite similarly, the objective for an optimizer is to replicate an image by only using brush strokes.

This is a similar task to what genetic algorithms already can do and have done in order to approximate images by other geometric shapes or even smaller photos (also known as the popular photo mosaic effect).

6.1 Motivation	15
6.2 Neural Renderer	17
Data Set	17
Architecture approach by Ha	22
Training	23
and Schmidhuber 'World'	23
Results	23
model	23
6.3 Stroke Approximation	23
Data Set	23
Weapon of Choice	24
Optimization Based Approach	28
Image Placing and Blending	32
Style Transfer	39
Optimization Details	39
Results	40



Figure 6.1: A typical set of brushes and spatulas used for oil paintings.

reformulate this

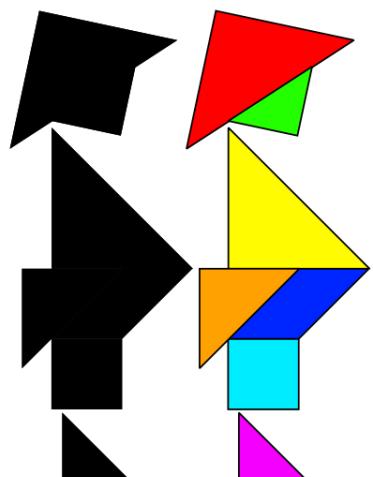


Figure 6.2: An Example of Tangram.



Figure 6.4: Photo mosaic of Starry Night using only images by the Hubble Space Telescope. http://www.astro.uvic.ca/~alexhp/new/figures/starrynight_HST.001.jpg



Figure 6.3: Starry Night approximated by a genetic algorithm using only circles. <https://effyfan.com/2018/03/02/w6-van-gogh-flowfield/>

Genetic algorithms follow a random sampling approach that 'evolves' like genomes do. Basically starting with a random set of circles that are parameterized by their position, radius and color; it then chooses the most successful samples and resamples again in a region around these. This process is repeated again and again, until a certain level of convergence is reached.

As well as this does work, it is very much computationally expensive as most samples will not fit the image, thus searching for the tiny set fitting shapes requires to evaluate all the bad shapes as well. Since brush strokes have many more degrees of freedom and artworks usually consist of upwards of a few thousand brush strokes this will not be applicable to this problem until computational resources have become a few magnitudes more powerful.

This premise can be overcome, though, by using something that is known as a differentiable renderer.

Differentiable renderers allow for the previously described task to be feasible as random sampling is replaced by gradient descent. A differentiable renderer is capable of creating shapes in the pixel domain by only using differentiable operations. Ordinary renderers do not have this property as they rely on faster operations that are not differentiable as most renderers are not used in this context. Also, creating a differentiable pipeline to render a circle from the tuple (position, radius, color) is a much harder task than calculating the shape of a circle beforehand and then projecting the shape onto the pixels in an image and coloring them accordingly. Nonetheless it is theoretically possible to create such a renderer.

Going back to the task of rendering brush strokes, it would be even harder to think of a pipeline that uses only differentiable functions and is able to draw brush strokes of a reasonable quality from a set of parameters.

This is where neural networks give an opportunity to avoid this problem. Neural network are inherently differentiable and previous works have shown that they are capable of high resolution and high quality conditional image generation. One could easily think of a generator setup that takes as an input the parameters of a brush stroke as well as some noise and then outputs an image of the according brush stroke with possibly some variability in outline shape or other characteristics.

This principle was proposed by [[japanese neural renderer](#)] as it facilitates training of reinforcement learning based networks.

Inspired by this, the approach becomes more clear. First, a differentiable renderer in form of a neural network is trained. Then this renderer is used by an optimization procedure that uses gradient descent to approximate an artwork as a input parameters batch of the renderer.

Both steps require some tricks to avoid pitfalls like computational limitation which are outlined in the following two sections.

6.2 Neural Renderer

The neural renderer is inspired by Nakano usage of a neural renderer to improve SPIRAL [1] as well as Huang et al.'s efforts on using deep reinforcement learning for drawing images with a differentiable renderer [2]. The renderer is required to be differentiable and should be based on a rather simple architecture. Especially since complicated architectures impose computational burdens and could possibly distort the gradients for optimization.

Data Set

Unfortunately there is no data set available for this task which means that the data set must be specifically created for this approach.

There are several sources for brush strokes that will be evaluated in the following part. The main focus will lie on three qualities for each data source:

1. Suitable data format. As brush strokes are meant to overlap later on, it data should already provide information about opacity or transparency, favorably in the common RGBA format, which has a fourth alpha channel to hold opacity information.
2. Size and versatility. Only with enough data available it will be possible to train a renderer reliably.
3. Image quality. Brush strokes should be as close to real world brush strokes as possible to ensure high quality renderings later on.

Brush Stroke Images

There are multiple sets of handrawn brush strokes available online. Most notably there is a set of various well classified colors and brush styles created by 'zolee' on the platform [onlygfx.com](#). It consists of approximately 1000 brush strokes that mostly follow rather straight

reference this

horizontal paths. Brush strokes are mostly grouped by color and painting technique (oil, acrylic, watercolor...). All images are in the PNG format and the area around the brush stroke was made transparent in a post-editing step.

This data set has the advantage that it consists of real world brush strokes that were painted under presumably reproducible conditions. On the other side, brush strokes are of mostly the same width throughout the data set and also do not come with information which path the brush took or any other non-visual information. Also, the data is very sparse. Many color shades are not represented which means that the generator would have to 'guess' them or simply would not be capable of rendering any brush strokes in this color.

It seems that this data set would be nice to replicate real world brush strokes as images but limitations to the data make it unlikely that a generator could learn a coherent representation from this.

Painting Libraries

The mentioned work of SPIRAL [3] relies on opposite data to real world images. They used the painting library 'libmypaint' [[libmypaint](#)] to generate brush strokes from parameters in real time during training.

The obvious advantage of this and other painting libraries is the fact that one can fully control the output through parameters. This makes it much easier as the whole space of input parameters for the renderer can be covered and avoids pitfalls like they were described in ??.

Still, this data set falls short regarding the authenticity of rendered strokes. Especially the inner area of the stroke shows a uniform color which is far from what real brush strokes would look like.

This data set is better suited for our task than the given images are but will tend to make all rendering look a bit 'cartoonish' or flat, which could in turn limit convergence during the latter optimization process.

Fluid Simulation

Fluid Paint is a project by David Li [[fluidpaint](#)] that uses simple fluid dynamics to give artificial brush strokes a more plastic look. It is implemented in JavaScript and OpenGL.

There is a C++ based version in the git-repository of SPIRAL which is also fitted with python bindings by Yaroslav Ganin.

Using these python bindings, it is possible to generate brush strokes locally outside of a web browser.

The quality and controllability of fluid paint falls right in between the two previously mentioned datasets. The generated brush strokes look distinctively better than those generated with 'libmypaint' but still lack the quality of the real world images. Concerning controllability, FluidPaint allows to control the path of the brush stroke handle rather than the brush stroke itself. This is a vast improvement over the images of 'zolee' but induces some offsets to a given path as opposed to 'libmypaint'.

It seems that this is reasonable compromise between the previously mentioned data sets. Although real time data generation in not possible with the library, it can be parallelized to allow for the creation of large data sets in a reasonable time frame.

Even though this data set still has some weaknesses, it comes in as the probably best choice for training a differentiable renderer because of the noted reasons.

Other honorable mentions are painting programs such as which allow for even more authentic brush strokes but lack a well documented interface in order to generate a vast number of brush strokes.

add similar samples for all brush stroke methods.

these two weird software stuff things

talk about RGBA advantages

Brush Stroke Formalism

With the means of data set production seized, what is left is to formulate the parameters that define the brush strokes. These parameters must quantify the following three properties of brush strokes:

- ▶ color
- ▶ thickness
- ▶ path

cut this joke

The easiest of these three properties is quantifying the color. Naturally, computer vision relies on the RGB format the defines color as a set of three 8-Bit integer values between 0 and 255. As for path and thickness these two properties depend on the given coordinate system. FluidPaint represents the canvas as a 2D plane in the $[0, 1]$ range, thus it makes sense to follow the same representation.

Thickness thus becomes a value in $[0, 1]$ for each brush stroke, where 0 is and infinitely small brush stroke and 1 is a brush stroke with the width of the canvas. As both the edge cases do not make sense, the range is constrained to $[.03, .2]$ which includes only brush strokes that are visible and also do not cover the whole canvas.

Quantifying the path now is a little more tricky. The fluid dynamics simulation that FluidPaint uses relies on internal time steps at which the equations are evaluated and subsequently rendered. At the same time each step allows only a linear motion of the brush handle between positions a and b . This means that any curved paths must be split into linear/straight segments, that should resemble a curved line. As more steps mean longer simulation times and fewer steps mean edgy movement, a value of 20 time steps per stroke is picked. This is equivalent to the number of steps that SPIRAL used in their implementation.

this is actually different as the steps depend on the lenght of the path. Look this up

Another problem becomes how to express a curved path in numbers. The easiest representation would be a sequence of points which make up the path of interest. This would give the highest versatility but at the same time introduce a noticeable amount of parameters as each point consists of 2 coordinates rendering 40 values. These values are also not independent but should follow a reasonable path as otherwise the resulting brush stroke would look rather like a random walk than a

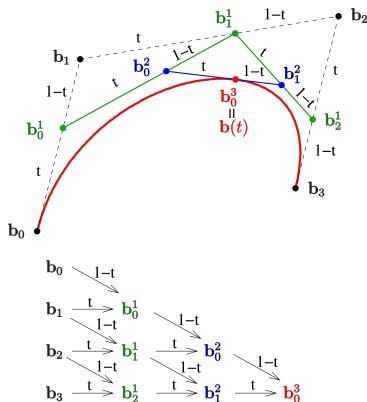


Figure 6.5: Sample of a 3rd degree Bezier curve, using the De-Casteljau-algorithm, <https://de.wikipedia.org/wiki/B%C3%A9zierkurve#/media/File:Bezier-cast-3.svg>

believable brush stroke. Since works such as SPIRAL or Learning to Paint face a quite similar obstacle, their solution should be applicable in this case as well. Both works use so called 'Bezier curves' which parametrize curved paths by a limited set of numbers.

Bezier curves were invented in order to parametrize curved paths not through checkpoints but analytical parametrizations for computer models. They can be of different orders which allows them to follow more complicated paths but in this case the simples form – the first order – is sufficient. It defines a curve through its start and end point as well as a control point. The curve is then defined as the path of a point over the time interval $[0, T]$. First one connect the start and end points with the control point to get two lines in return. Along the lines on defines two points that move in a linear fashion along their lines within the time interval $[0, T]$. Then these two points are connected in the same way by a third line, which again has a point moving along its direction over the span of $[0, T]$. As the first two points that define the third line will move, the lines orientation will change as well, thus translating the linear movement of the point into a complex curved path. The path that this point then takes in $[0, T]$ then defines the Bezier curve. A first order Bezier curve will only bend into one direction or follow a straight path. For higher orders the displayed process can be applied iteratively and allows for more complex curves but as brush strokes usually follow a quite simple path and fewer parameters are preferred, Bezier curves of first order are chosen as parametrization.

Ultimately this gives 10 values that are sufficient to parametrize brush strokes with certain constraints:

- ▶ 3 two dimensional coordinates that define the Bezier curves (6 values).
- ▶ 1 thickness parameter.
- ▶ 3 values in RGB space.

Data Constraints

Given the parameters listed in section 6.2, the data still needs further constraints to facilitate training of the generator even further.

Section 6.2 already hinted at the impracticality of online data generation. A rough estimation by timing the rendering of 100.000 FluidPaint brush strokes reveals that a dedicated CPU server is capable of generating 300 strokes per second when using all of its resources. Keeping this in mind, a neural network with batch size 32 is limited to ≈ 10 iterations per second which would mean a clear bottleneck. Thus, it seems advisable to generate data beforehand with enough samples to cover the data space sufficiently. This will allow for much faster access to data, as individual data samples are relatively small and can be stored in a binary data file such as HDF5.

Besides this constraint to the amount of data available, another set of constraints will be introduced to reduce the data space to 'valid' brush strokes only. 'Valid' brush strokes will be defined as brush strokes which resemble real world brush strokes to a certain degree. This primarily concerns two relations within a brush stroke:

experiment in appendix

- Its width-to-length ratio.
- Its curvature.

The width-to-length ratio will be restricted to brush strokes that are at least two times as long as they are wide.

$$\|\underline{s} - \underline{e}\| \stackrel{!}{\leq} 2 \times (\text{brush size}) \quad (6.1)$$

Due to the simulation background of FluidPaint shorter brush strokes will show some artifacts due to the bristles' length in the simulation which depends on the width of the stroke. Another reason for this, is the intended use case which will focus on van Gogh paintings. As van Gogh did not practice pointillism most of his strokes have length to them, which brings such a constraint in line with some characteristics of van Gogh's style.

The same argumentation can be done for the curvature: Most brush strokes (especially those by van Gogh) have a certain 'flow' or 'smoothness' to them, which can be described as by using strokes with large radii of their curvature and without any corners in a strokes' path. Thus data set will also be restricted to strokes which follow these descriptions. In order to achieve this with random sampling in mind, a multivariate gaussian distribution is placed between start point (\underline{s}) and end point (\underline{e}). The two axes are rotated such that the short axis is in line with the vector $\underline{a} = \underline{s} - \underline{e}$ while the other sits orthogonal. Then both axes are scaled with $\|\underline{a}\|_2$ and also the handpicked values $\frac{1}{200}$ and $\frac{1}{25}$ for along \underline{a} and orthogonal to it, respectively. Figure ?? shows samples from this distribution for an exemplary brush stroke. This distribution is intended to follow that of brush strokes as they would appear in the real world. The majority of brush strokes will be straight or just slightly bent due to the maximum of the PDF being at the center of \underline{s} and \underline{e} . Bent brush strokes will mostly be symmetric as the long axis of the multivariate gaussian is orthogonal to \underline{a} . Still, there will be strokes that have their bent towards either end of the brush stroke as well as some strokes with a high curvature. The area of interest, though, will be densely populated as intended.

draw figure in matplotlib or plotly or something

$$p(\underline{c}|\underline{s}, \underline{e}) = \mathcal{N}(\mu, \Sigma) \quad (6.2)$$

$$\mu = \frac{\underline{s} - \underline{e}}{2} + \underline{e} \quad (6.3)$$

$$\Sigma = \begin{pmatrix} a_x & 0 \\ 0 & a_y \end{pmatrix} \quad (6.4)$$

$$\underline{a} = \underline{s} - \underline{e} \quad (6.5)$$

The color of the brush strokes is not constrained as the color distribution of the target data set is not known at this point.

why not van gogh color distribution?

Data Set Creation

The data set will be created with 100.000 samples that follow the constraints that were presented in section 6.2. As underlying distribution the uniform distribution is chosen as it allows a more evenly coverage of the data space.

First, a set of start and end point as well as brush size is drawn and checked against (6.1). If the constraint is not met, the set will be redrawn entirely. In case the constraint is satisfied, a checkpoint is sampled according to (6.2). If \underline{c} lies outside the render window, the checkpoint will be resampled. At last a color in form of an RGB value is sampled from a uniform distribution as well.

The resulting tuple of start, end and control point, brush size and RGB color is then added to the data set. Before rendering starts the values of \underline{s} , \underline{e} and \underline{c} are scaled with the hand picked factor of 0.7 to ensure the brush strokes are rendered completely and not cut by the edge of the render window size. At last the brush strokes are rendered according to the data set and added as well.

The render canvas size was chosen to be 64x64 pixels for several reasons: First, even with such a small canvas size, training for the renderer takes about one day. Secondly, the larger the render canvas size becomes, the deeper the renderer needs to be which results in more computational overhead in the optimization routine as well as more layer through which the gradient has to be propagated. And lastly, as there will be upwards of a thousand brush strokes in a single image, increasing the canvas size to 128x128 would require four times as much memory per rendered image. As 1000 brush strokes would already account for $1000 \times 64 \times 64 \times 4B \approx 16.4\text{GiB}$ a fourfold increase would be significant.

As a last step, the data set is renormalized to the range $[-1, 1]$ for convenience and to facilitate training as well.

Architecture

The architecture of the brush stroke generator follows that of an inverse VGG network. It is widely used and has shown in previous works that it should be capable of handling this task. The architecture consists of three dense layers at the beginning, followed by a 2 times upsampling layer as well as three convolutional layers. The same set of a 2 times upsampling layer and three convolutional layers is repeated until the target size is reached. After the last convolutional layer a tangens hyperbolicus function is applied, to restrict the output to the $[-1, 1]$ range. As part of the hyper-parameter search different tweaks to the architecture have been tested:

- ▶ An additional noise input at every layer with a size equal to that of the existing signal.
- ▶ Additional information about the position in the pixel grid in every layer, so called CoordConv [[coordconv](#)].
- ▶ Various combinations of activation and normalization functions.

The discriminator is designed after the same principles and resembles a VGG encoder network. First three convolutional layers are applied, followed by a downsampling or pooling layer. This is repeated until a target resolution of say 4x4 pixels is reached. Then a set of three dense layers is applied to give one final prediction per sample.

[check the details](#)

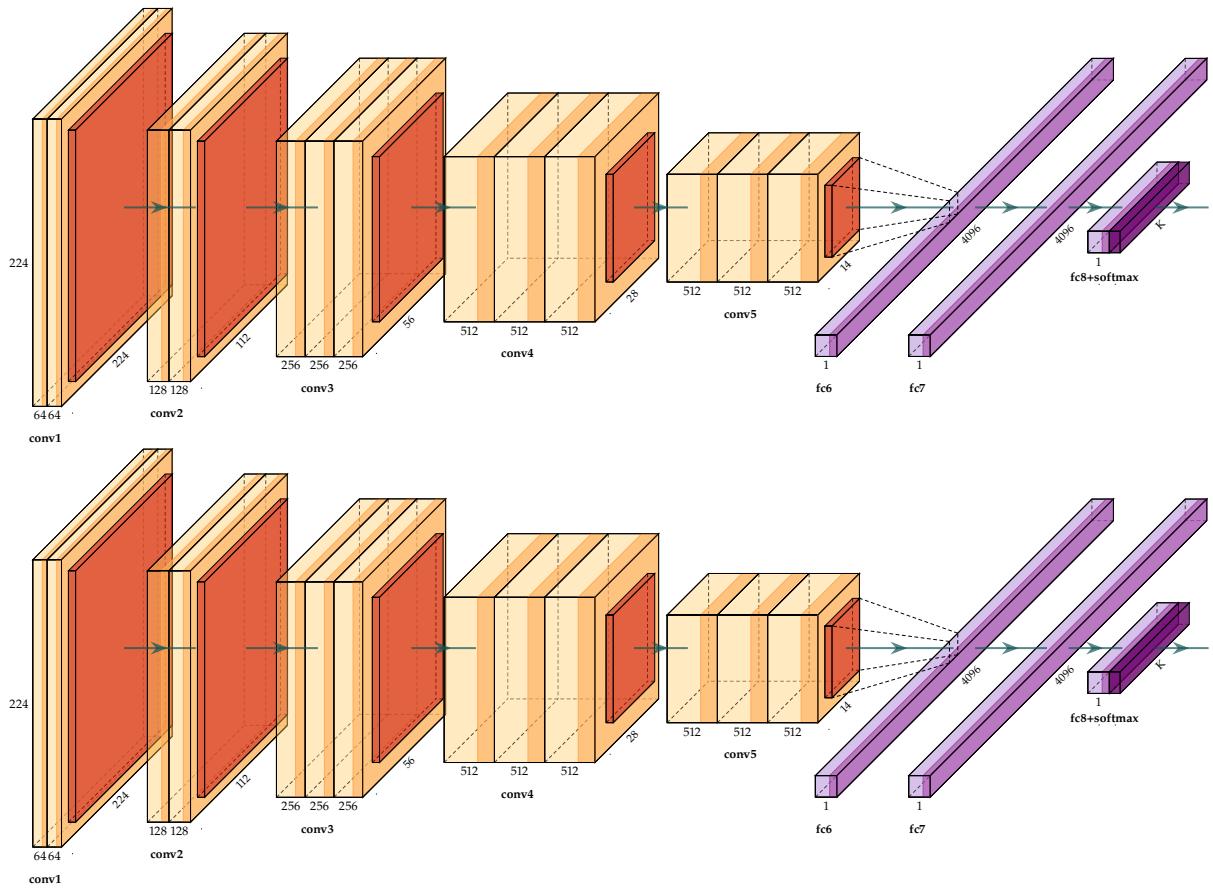


Figure 6.6: Visualization of the generator and discriminator architectures

Training

During training the L2 distance and the FID score were evaluated as metrics. The FID score became necessary as the visual comparison of the generated samples has proven difficult between different runs. The L2 distance did not qualify as a sufficient metric for later training stages as the stochastic nature of the brush strokes puts a threshold on how short the L2 distance can become.

To stabilize training further a two-time-step update rule was implemented.

write this down and look at the tricks that were used

Results

pick some results and present whether they are any good

6.3 Stroke Approximation

Data Set

First off, for this section the data set will be presented. This is due to the data defining some demands for the networks later on. The data set for the optimization task should also meet a few requirements.

In order to be able to focus on the brush strokes in an image the data set should consist of relatively high resolution images. This becomes necessary as most brush strokes should fit into a 64x64 window any larger strokes could not be approximated with a single renderer and thus would falsely be constructed of multiple brush strokes.

This requires more information per image than simply the resolution as the scale of the image plays into this requirement as well. Ideally each image should be accompanied by its size as this allows to rescale all images accordingly, given that one knows how large a typical brush stroke is.

At last the painting's technique must be oil on canvas or similar techniques as this is what the renderer has been trained for.

All these requirements on the data can be met by using data directly from the Van Gogh Museum in Amsterdam that is freely available online. Each high resolution image is categorized by its technique as well as the period in which it was painted accompanied by information on the measurements of the image.

Weapon of Choice

The next stage of training takes the now pre-trained neural renderer and combines it with the task of approximating brush strokes in images. In theory the neural renderer should allow to gather gradients in parameter space even though meaningful losses are calculated in image space, as explained in section 6.1. There, it was outlined already that the weapon of choice will be an optimization procedure that relies on the capabilities of the neural renderer.

But this is not the only possible approach to dissecting images into sets of parameters. As explained in 5.2 there exist a variety of approaches for this kind of task. Some of these promise fast generation of parameters from images which makes the optimization approach in this work seem like a step back at first. The following section is dedicated to justifying the chosen approach by comparing state of the art implementations of existing approaches.

The targeted task is to approximate the representation of an image with roughly 1MP by the means of ≈ 10.000 brush strokes (see section ??).

Genetic Algorithms

Genetic algorithms are possibly the simplest approaches. As it was laid out in sections 6.1 and ??, genetic algorithms use random sampling and previous best results to approximate fitting solutions to the rendering problem which was described in section ??.

Current state of the art solutions are capable of finding a solution in $\approx 1h$, when searching for an approximation of a 1MP image with simple geometric shapes like circles or triangles. This accounts for roughly 1×10^9 sampling steps. This time frame varies depending on target shape density, target accuracy, sampling runs per shape and degrees of freedom (with the latter requiring more sampling runs).

Looking back at section ?? each brush strokes has 10 degrees of freedom with some inter-dependencies between them. Also it is known that the obtained neural renderer is capable of generating $\approx 300 \frac{\text{images}}{\text{s}}$. This means that rendering the same amount of sampled images will take:

$$\frac{10^9 \text{ samples}}{300 \frac{\text{images}}{\text{s}}} \approx 3.33 \times 10^6 \text{ s} \approx 926 \text{ h} \approx 38.6 \text{ d}$$

Which is an impossible amount of time to spend **per image**.

At this point the number of samples has not been corrected for by the higher number of degrees of freedom of this problem.

Ultimately this means that genetic algorithms are not an option for this task.

Brush Stroke Extraction

Next in the line of approaches are algorithm based approaches that use standard computer vision techniques to extract brush strokes from an image and parametrize them.

Besides texton based image characterization [**textons**] and pure filter based approaches [**filters**] there have been approaches to extract brush strokes or some of their characteristics [**brushstrokecharacteristics**] [**brushstrokeextraction**].

The latter would pose a valid option for the main goal of this thesis if it were to characterize brush strokes reliably. Unfortunately the best existing techniques fail to detect brush strokes equally over the whole image but only identify the most significant ones (see figure ??). Other approaches are only able to extract only few characteristics like the orientation of a brush stroke, which proves insufficient as well.

show some of the extracted
brush strokes

Stroke Based Rendering

A different field that also uses algorithm to approximate an image through brush strokes and also wants to achieve some stylization along the way is Stroke Based Rendering or Painterly Rendering. While early works relied on interactive approaches, later publications then were able to fully automate this process. Judging from tests by the authors themselves such an approach would take between and hours per image to fully render it this way. As figures ?? and ?? show, these approaches obtain similar results to genetic algorithms and tend to draw an image from a coarse scale to a finer scale instead of locally coherent.

this number

this number

At last these approaches were not intended to be used to obtain brush strokes from a painting but rather focus on stylizing images. This renders such an approach unfit for the goal of this thesis as well.

which one was first?

Drawing Networks

Lastly, drawing networks are the newest iteration of approaches in this field. Beginning with there have been approaches the make use of feed-forward or recurrent neural networks combined with either supervised training or deep reinforcement learning. The best results by these approaches are shown in figure ??.

Noticeably, all of these images have a maximum resolution of 256 pixels along the longer edge. There have not been any approaches yet, which are able to go significantly beyond this limitation. Especially, the high computational costs of training recurrent neural networks seems to be an obstacle when shooting for higher resolutions. As it will be outlined in more detail in 6.3 such resolutions can not be deemed sufficient when looking at individual brush strokes.

This, also rules out drawing networks.

Combined Approach

Even though there have been plenty of previous approaches to the task of extracting brush strokes from images or to similar task like rendering images through brush strokes; none of these quite meet all the requirements that have been posed for this task. Namely, high resolution input images, brush stroke focussed rendering and/or retrieval, limited resources and realistic depiction of brush strokes.

Even though brush stroke extraction seems to be most fit for the main asset of this thesis, stroke based rendering and drawing networks both bring features into the mix, that seem capable of making the task at hand more feasible. Drawing networks for once introduced differentiable renderers as a tool to facilitate training. Stroke based renderers achieve a parametrization implicitly by focussing on replicating the entire image instead of extracting brush stroke by brush stroke.

This can be combined into a unified approach that uses the differentiable renderer and the objective of recreating an entire image. The resulting approach does not use re-sampling like genetic algorithms do, but can rely on gradient descent to converge to a solution significantly faster.

The limited capabilities of such a renderer would also guarantee that any approximation is composed of only valid brush strokes instead of single pixels it would be the case with normal generative models.

All in all such an approach fits best to the goal of this thesis as it will generate an approximation of a target image that ideally resembles the target as closely as possible while being limited to the use of brush strokes.

The decision of whether to use an optimization based approach is also linked to another question: Whether to use a sequential or parallel placement of brush strokes.

As an example most drawing networks rely on a sequential approach as do some stroke based renderers. Intuitively it also makes sense to use a sequential approach as artists also place their brush strokes sequentially on the canvas. Yet there is a major difference in how existing computer

vision based approaches place their strokes compared to artists. Due to the loss functions that are used when training drawing networks or stroke based renderers the resulting images tend to be made up entirely of large canvas-filling brush strokes which reduces the L2 loss very fast.

This is contrary to an artist which would maybe start with some sketches first but then fills the canvas not with a brush that is roughly the size of the canvas – as these algorithms do – but with a significantly smaller brush that requires many individual strokes (often with some pattern) to fill large areas. Thus, sequential approaches do actually differ significantly from how an artist would paint.

In contrast to this parallel approaches which predict the whole painting in one step are not wide spread due to the computational pitfalls that come with predicting brush strokes for a whole painting at once compared to predicting only a few at a time. Also, parallel approaches require to deal explicitly with the order in which brush strokes should be placed on the canvas, as it will be explained in section 6.3. Still, if one considers the optimization procedure instead of a policy network, it becomes clear that there are interactions between strokes that are locally close. An example would be one stroke in the foreground changing its path and revealing the canvas beneath, then another stroke in the background should cover this up if the color matches.

As usually artists plan their future brush strokes when painting the background it would require a drawing network to plan far ahead which is impossible to do with an optimization based approach.

Another argument for parallel optimization is the actual reliance on actual visual proof for approximating brush strokes. Sequential approaches will again and again cover up previous brush strokes until maybe only a few pixels of the background brush stroke are visible. Even if these few pixels match the target image well it will be a shot in the dark whether this is really how the artist has drawn the background, which in return introduces a great deal of noise to the result. A parallel approach on the other hand should only place brush strokes where there is visual evidence that there is a brush stroke and concentrate less on possible background arrangements.

Pitfalls of Feed-Forward Approaches

In the course of this thesis there were experiments targeting a feed forward approach before ultimately tending toward the optimization based approach that is presented.

The main reason for this change in direction can be pinpointed to two problems that emerged. The computational burden that of a feed-forward approach can be approximated when looking at existing feed-forward approaches of drawing networks. Without compromising resolution and this images quality significantly, training a network can hardly be realized. Still it was possible to show a basic implementation of this for very simple data like the cMNIST data set.

Another problem that occurred was the placement of brush strokes on the canvas. As artists are not bound to the same pixel grid as images are, they can place brush strokes freely on the canvas. More so, they can pack

margin explain cmnist

add image that were drawn by ff network

brush strokes densely in one area while distributing them broadly in another.

As most neural networks in computer vision are CNNs, they are not able to allow for a similar behaviour as they will always follow a grid layout of various resolutions.

Repeated experiments with either displaceable grid cells or stacked signals have proven too complicated to manage for a convolutional network architecture and also seemed to scale badly when implemented in fully convolutional manner.

Thus an optimization based approximation was chosen as it offers good approximations at high resolutions with manageable computational overhead.

Optimization Based Approach

Since the previous section ?? already specified why an optimization based approach has been chosen over a feed forward or recurrent approach, this section is meant to explain a little more about the optimization procedure itself.

Rendering Layout

Fundamentally, the optimization procedure is inspired by stroke based rendering procedures. Also, it could be compared to the style transfer approach by Gatys *et al.* [gatys], where, instead of pixels there are parameters optimized. The difference to normal stroke based renderers, though, is the limited size of a rendered brush stroke in this work's renderer.

This poses a significant challenge that might not be obvious at first.

Ideally the optimization procedure should be able to place strokes freely on canvas, as this allows for an unbiased approximation. Furthermore, this allows to allocate many small strokes in areas where the artist placed many strokes as well and use fewer and wider strokes in other areas during the optimization procedure. But due to various limitations which were explained in ??, the renderer is not able to render single brush strokes in a 1MP frame. Similar approaches only perform on relatively small canvas sizes, likely due to this issue.

add ref

One work-around for this would be to dissect the target image into many smaller patches and then run the optimization procedure on these individual patches. This could be compared to a grid of renders where each grid cell is the center of many renderings at the same time. Right away, a major issue with this becomes obvious, which are the edges where the grid cells are joined.

add figure to this

As a grid structure will almost always differ from the inherent distribution of brush strokes in an image, grid edges will more often than not, separate brush strokes between two grid cells. A simple solution to this problem is going from a stacked grid structure to an overlapping grid.

graphics!!!!

This hides obvious edges between render windows as every edge coincides with the center of a different render window. It is easily realized by choosing a lattice vector size smaller than the render window's dimensions.

Still, this kind of initialization requires a very even distribution of brush strokes with only small dislocations at each point in the lattice. As this is not the case and stroke densities will vary locally, the grid layout is prone to lead to enforcing a grid-like layout of strokes were there isn't any. Mainly due to the inability of the grid to account for local changes in density and the following propagation of error. This would start with a single region of high density strokes in the vicinity of one grid cell which would ideally result in a narrow stroke to achieve high accuracy. Neighboring cells then have to shift their strokes towards the center of that grid cell to account for the free space that is not covered by the narrow brush strokes. This shift must then be accounted for by the next neighboring cells and so on, which will cause all strokes in a raw or column to shift towards this one spot with a high stroke density.

Now, a painting usually has many such high density areas which would require the optimization procedure to balance the shift that is caused by these areas. As a result the strokes are likely to not shift at all as shifting will equal out for many such areas of high density. Subsequently, an area of high stroke density will be not have enough strokes available in its local region and thus will be covered up by a single broad stroke as this minimizes the L2 loss.

The core of the problem is the previously imposed lattice structure that propagates local density shifts along its principle axis.

One possible solution to this is getting rid of the lattice structure and replacing it with a more random structure that also covers the image sufficiently. This can be accomplished by using **super-pixels [superpixels]**. Super pixels come from an earlier era of computer vision and are often used in image segmentation tasks (**[img segmentation with SP]**) but super-pixels are also a popular starting point for brush stroke extraction algorithms **[brushes stroke extraction]**.

Basically, super-pixels are disjoint groups of pixels in an image that would usually combine pixels with similar colors in a local region. Straight away, it is obvious why this is interesting for brush stroke extraction. In theory the distribution of super-pixels will not follow a grid-layout as the previous approach and the location of super-pixels should relate to the given color distribution in the image. It is easy to imagine that the location of the super-pixel centers would be a good prior of locations of render windows as well. Also, as the colors of pixels inside a super-pixel should be similar, one can use the mean color of a super-pixel as an initialization for the color of the brush stroke.

Ultimately, a super-pixel segmentation will be used to infer positions for render windows as well as the color initialization of each strokes.

Rendering Order

Another problem that will come up during the optimization procedure is the order in which strokes are rendered. Real world brush strokes are

also subject to the same issue as later brush strokes will always be placed on top of earlier brush strokes, with no way of changing this (see section 6.3).

As the optimization based approach relies on parallel optimization of brush strokes, it must be decided which strokes are in the foreground and which are in the background. Otherwise, as this would randomly change, edges in the image might be obstructed and optimization could oscillate between solutions where different strokes lie in the foreground. It could also lead to brush strokes not overlapping but covering disjoint areas. All of these outcomes would be unfavorable as it tends to produce worse results in the end.

The solution which is presented in this thesis is an additional parameter that describes a brush stroke's accuracy. The accuracy is defined as the L2 distance of each stroke's pixel to the corresponding pixel in the target image multiplied by this each pixel's alpha value. This removes any pixels which are of no interest from the loss and focusses only on the rendered pixels.

$$\text{accuracy} = 1 - \frac{1}{N} \sum_{p \in \text{pixels}} \|p - p'\|_2^2 \text{ with } p = p' \quad (6.6)$$

improve this

The resulting value describes how well the pixels of the rendered stroke match their respective pixels in the target image. Consequently, any brush stroke with a higher accuracy will be more faithful to the target image than strokes with a lower accuracy. Placing these brush strokes in the foreground should thus result in a smaller L2 loss than the other way around. Vice versa, brush strokes which connect two same colored areas will aggregate a lower accuracy as the brush stroke is compared at the intersection as well. A rendered brush stroke which fits to the foreground brush stroke will not be affected by such a penalty thus getting a higher accuracy and laying on top of the other brush stroke.

Notably, the accuracy should not be included in the brush stroke's loss, as this would prohibit background strokes from covering larger areas and result in behaviour that is similar to non-overlapping issue previously described. Thus, the accuracy of each stroke will be calculated as it is rendered.

Initialization

The following section will focus on initialization details for all parameters of a brush stroke, their position and the confidence value.

Besides the original 10 parameters of each brush stroke, which were explained in section 6.2, the previous two section introduced an accuracy parameter for ordering and two translation parameters which define the position of the render window along each axis. All of these parameters must be initialized before the optimization procedure starts. Ideally, the initialization should not introduce any bias to the optimization process. At the same time an initialization should facilitate training and accelerate convergence in early stages of optimization.

Unfortunately, the placement of the render window will surely enforce a bias on the optimization as hinted in section 6.3. There a super-pixel initialization was motivated and for the translation parameters as well as the color of the brush strokes. Subsequently, the translation parameter for each render window will be equal to the position of the weighted mass center of its respective super-pixel.

The initial color will be taken from the mean color value over the super pixel.

The brush size will be initialized with the minimum possible value. This will let brush strokes not overlap at the beginning of optimization. Only when the brush strokes already roughly fit their local region they shall intersect and be ordered by their accuracy. Therefore, the initial accuracy will be 0 everywhere, as the accuracy is recalculated after every optimization step.

Other patch parameter, notably s , e and c , will be initialised using a narrow normal distribution with $\sigma = .1$ and values clipped to $[-1, 1]$. This is done, as there is no prior information available on how the brush strokes are oriented. Instead, this approach relies on the optimization procedure to be minimally biased by this initialization of the path variables.

Partial Updates

One major problem that emerges when building an optimizer based on the information provided up until now, is the limitations of GPU memory. Even as it might not be obvious at first, the optimization procedure imposes a huge requirement for memory on the graphics card. This is mainly due to two parts of the training:

First, the number of brush strokes can easily become very large, if large enough images are used as input. As the render window size of 64x64 already suggests the brush strokes are relatively small and mentioned in section 6.1, a painting can consist of a few thousand brush strokes. This would equate to a batch size of a few thousand for the brush stroke renderer. Tests have shown that on an NVIDIA GTX 2080 Ti with 12GB GDDR5 VRAM the maximum number of brush strokes rendered in parallel is ≈ 256 . Obviously this is one to two orders of magnitude smaller than what would be needed to optimize all stroke in the painting in parallel. Still, there is a way around this bottleneck, by optimizing the image not as a whole but as smaller patches consisting of 256 brush strokes at a time, giving a partial update routine.

Each patch comprises the 256 nearest render windows to a randomly sampled location on the canvas. These 256 brush strokes are then rendered from their parameters in order to obtain a gradient later on. Then the strokes are placed on the canvas or ‘padded’ and then blended or ‘stitched’ after they were ordered according to their accuracy. At last, the loss is calculated and backpropagated to update the parameters of the patch.

As this will not make a difference for any brush strokes at the center of the patch, border brush strokes (those which do not solely neighbor brush strokes that are also optimized) will be affected by this. This is the reason why it makes sense to save the rendered brush strokes for all parameters at any time forming a **render image catalogue**. Equally the

collection of brush strokes that shall later compose the image shall be called a **parameter catalogue**. With the render image catalogue at hand, it is possible to use the previously rendered brush strokes to render the image as whole with the newly rendered brush strokes of the image patch embedded.

As this will cause the border strokes to be surrounded by other brush strokes to all sides (even if not all of them are freshly rendered) the effect of the partial update routine vanishes.

The other problem that will occur for very large input images, is that the stitching of brush strokes itself takes up a huge mount of RAM. Going into detail, each brush stroke must be placed on the virtual canvas individually, where the canvas' size is that of the input image. This would equate to a couple thousand 1MP images being stored in RAM before they are stitched to a single 1MP images. As a single 1MP image carries roughly $1000 \times 1000 \times 4\text{channels} \times 8\text{Bit} = 32 \times 10^6\text{Bit} = 4 \times 10^6\text{Byte} = 4\text{MB}$ of information, a few thousand of these will easily eat up the RAM of most graphics cards.

Luckily, the previous work-around for optimizing only 256 brush strokes will work as well without rendering the full image at every step. Since most brush strokes are not re-rendered anyway and thus will not be supplied with a gradient, their main task is to regulate losses for edge strokes of the image patch. As this does not need far away strokes but only those close the strokes that are optimized, a ring of pre-rendered strokes around the re-rendered patch will suffice.

approximate the amount of brush strokes needed to surround the patch

This allows to reduce the amount of involved brush strokes per optimization step from a few thousand down to a couple hundred. Besides allowing for the partial update routine to be performed at all, it also should increase performance significantly compared to an approach that involves all brush strokes at any time.

Image Placing and Blending

As the update and optimization procedure has now been explained thoroughly, it is now time to explain the process of placing and blending a rendered brush stroke a bit further.

After each stroke has been rendered it needs to be placed according to the translation parameters that were introduced in section 6.3. This requires dynamically placing each brush stroke inside zero-filled tensor.

By calculating the global position of each pixel in the rendered image individually, it is possible to scatter the pixels of the original rendered image into the larger zero-filled tensor and obtain a globally placed brush stroke.

What is more complicated though is the task of blending the resulting canvas-sized renders together or stitching them. Due to the canvas' alpha channel it is possible to blend only relevant information while the rest of the image will be ignored.

As far as conventional alpha blending goes, two images are blended by multiplying each pixel value with the alpha value of the top-layer image

while the background image is multiplied with the complement to the alpha value:

$$p_{x,y} = p_{x,y}^{\text{top}} \times \alpha_{x,y}^{\text{top}} + p_{x,y}^{\text{bottom}} \times (1 - \alpha_{x,y}^{\text{top}}) \forall (x, y) \in \mathcal{D}(\text{image}) \quad (6.7)$$

correct this equation and make it nicer

For multiple layers this process can be repeated in various fashions, after the strokes are ordered according to their accuracy. Either one could start from the bottom and blend the two back-most strokes, followed by the next third last strokes and so one, or one could start this process from the front with the two strokes in the very front being blended at first, then the stroke with the third highest accuracy etc..

Both of the approaches would yield the same result but differ only in the order in which they were blended. Subsequently, both methods will have $(n - 1)$ blending operations to compute per pixel.

By moving from the linear approaches to a tree-based approach this number will not decrease any further.

What makes it possible to reduce this number, though, is blending the strokes in a content aware fashion. As looking at figure ?? shows, the majority of pixels for each padded brush stroke is non-informative, as the alpha value is zero. This hints at possibility to turn the process of blending around and instead of merging layers subsequently as a whole, merging all layers at the same time by picking relevant layers for each pixel. Then all non-zero layers are merged in the same way as ordinary alpha blending works but with the number of layers being reduced and information inside a layer not being coherent anymore.

The easiest way to accomplish this is to first find the maximum depth over all pixels, where the depth k is the number of layers where the alpha value is not zero.

$$k = \arg \max_{p \in \text{pixels}} \sum_{i \in \# \text{layers}} \mathbb{1}(\alpha_i > 0) \quad (6.8)$$

Then the top k layer indices for each pixel are picked, which reduces the number of blending operations from $(n - 1)$ to $(k - 1)$. Importantly, the top k indices should not be ordered by their alpha values but remain in the order that was imposed by sorting according to the accuracy value. Otherwise, the order will most certainly be mixed up and the pixel with the highest alpha value will always lie on top instead of the pixel that belongs to the most accurate brush stroke. Especially as brush stroke renderings fade out towards their edges, this makes a significant difference.

Another way of accelerating the process of alpha-blending is vectorizing the process instead of iteratively applying the computations. To achieve this it is necessary to construct a tensor with the following properties:

For $\mathbf{I} \in [0, 1]^{H \times W \times 4}$ the image target, the shape will be defined as $S(\mathbf{I}) = (H, W, 4)$. Each alpha channel will have the values $\alpha^{hw} \in [0, 1]$ for $h = 0, \dots, H$ and $w = 0, \dots, W$.

The set of rendered and padded brush strokes \mathbf{J} will have the shape $(N, H, W, 4)$ with N depicting the number of brush strokes that ought to be stitched simultaneously.

Now, looking at each individual pixel in \mathbf{J} , which is described by $(z_n^{hw}, \alpha_n^{hw})$ for $n = 1, \dots, N$ and $z_n^{hw} \in [0, 1]^3$, z^{hw} describes the RGB values and α^{hw} the alpha-channel for a pixel at (h, w) .

A blending operation can then be defined by

$$z'^{hw} = \tilde{\alpha}^{hw} \cdot z^{hw} \quad (6.9)$$

$$\text{or} \quad (6.10)$$

$$z'^{hw} = \sum_{n=1}^N \tilde{\alpha}_n^{hw} z_n^{hw} \quad (6.11)$$

$$(6.12)$$

with z'^{hw} the resulting RGB values of the blended pixel and $\tilde{\alpha}^{hw}$ a vector that holds the merged alpha values for each pixel:

$$\tilde{\alpha}^{hw} = \begin{pmatrix} \alpha_1^{hw} & (1 - \alpha_1^{hw}) \\ \alpha_2^{hw} & (1 - \alpha_2^{hw}) \\ \alpha_3^{hw} & (1 - \alpha_3^{hw}) \\ \vdots & \vdots \end{pmatrix} \quad (6.13)$$

$$= \alpha^{hw} \odot \begin{pmatrix} 1 \\ (1 - \alpha_1^{hw}) \\ (1 - \alpha_2^{hw}) \\ \vdots \\ (1 - \alpha_N^{hw}) \end{pmatrix} \quad (6.14)$$

$$\rightarrow \tilde{\alpha}_n^{hw} = \alpha_n^{hw} \prod_{i=1}^{n-1} (1 - \alpha_i^{hw}) \quad (6.15)$$

Where \odot describes the element-wise product.

What is left, is to find a way to construct $\tilde{\alpha}^{hw}$ from α^{hw} .

For this an auxiliary matrix β^{hw} is constructed:

$$\beta^{hw} = \alpha^{hw} \times \mathbb{1}_{1 \times N} = \begin{pmatrix} \alpha_1^{hw} & \alpha_2^{hw} & \dots & \alpha_N^{hw} \\ \alpha_1^{hw} & \alpha_2^{hw} & \dots & \alpha_N^{hw} \\ \vdots & \vdots & \ddots & \vdots \\ \alpha_1^{hw} & \alpha_2^{hw} & \dots & \alpha_N^{hw} \end{pmatrix} \quad (6.16)$$

with

$$\mathbb{1}_{1 \times N} = \begin{pmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{pmatrix}^T$$

Then β^{hw} is strictly triangulated such that:

$$\gamma^{hw} = \beta^{hw} \odot \begin{pmatrix} 0 & 0 & 0 & \dots & 0 \\ 1 & 0 & 0 & \dots & 0 \\ 1 & 1 & 0 & \dots & 0 \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ 1 & 1 & \dots & 1 & 0 \end{pmatrix} \quad (6.17)$$

$$= \begin{pmatrix} 0 & 0 & 0 & \dots & 0 \\ \alpha_1^{hw} & 0 & 0 & \dots & 0 \\ \alpha_1^{hw} & \alpha_2^{hw} & 0 & \dots & 0 \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ \alpha_1^{hw} & \alpha_2^{hw} & \dots & \alpha_{n-1}^{hw} & 0 \end{pmatrix} \quad (6.18)$$

$$\rightarrow \delta^{hw} = 1 - \gamma^{hw} = \begin{pmatrix} 1 & 1 & 1 & \dots & 1 \\ (1 - \alpha_1^{hw}) & 1 & 1 & \dots & 1 \\ (1 - \alpha_1^{hw}) & (1 - \alpha_2^{hw}) & 1 & \dots & 1 \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ (1 - \alpha_1^{hw}) & (1 - \alpha_2^{hw}) & \dots & (1 - \alpha_{n-1}^{hw}) & 1 \end{pmatrix} \quad (6.19)$$

By multiplying the elements along each row in δ^{hw} one gets:

$$\epsilon_i^{hw} = \prod_{j=1}^N \delta_{ij}^{hw} = \begin{pmatrix} 1 \\ (1 - \alpha_1^{hw}) \\ \vdots \\ \prod_{j=1}^{N-1} (1 - \alpha_j^{hw}) \end{pmatrix} \quad (6.20)$$

$$\rightarrow \tilde{\alpha}^{hw} = \epsilon^{hw} \odot \alpha^{hw} = \begin{pmatrix} \alpha_1^{hw} \\ \alpha_2^{hw} & (1 - \alpha_1^{hw}) \\ \alpha_3^{hw} & (1 - \alpha_2^{hw}) & (1 - \alpha_1^{hw}) \\ \vdots \\ \alpha_N^{hw} & \prod_{j=1}^{N-1} (1 - \alpha_j^{hw}) \end{pmatrix} \quad (6.21)$$

This vectorized version of alpha blending will introduce a new possible bottleneck as it is, since β^{hw} will be a tensor of shape (N, N, H, W) , which will equate to

$$256 \times 256 \times 256 \times 256 \times 4B = 2^{36}B = 64GiB$$

alone.

This is where the previous content aware alpha blending tricks becomes useful. By computing β^{hw} only through the top k values of α^{hw} instead of the full tensor α^{hw} , the size will be reduced to

$$k \times k \times 256 \times 256 \times 4B = k^2 2^{22}B = k^2 \times 4MiB$$

As the shape is reduced to (k, k, H, W) .

Ultimately, this accelerates optimization by a factor of 2 – 3 as it will be shown in ??.

It must be mentioned that the upper boundary for computational complexity in using this kind of alpha-blending is $\mathcal{O}(N \log k)$ as the top k search is bound by this complexity.

calcualte this

Losses

In this section the different kinds of losses for the optimization procedure will be discussed.

First off, the L2 loss or **mean squared error** is an obvious choice for this task.

$$L_{\text{MSE}} = \frac{1}{HW} \sum_{p \in \text{pixels}} \|z(p) - z'(p)\|_2^2 \quad (6.22)$$

Since the MSE loss focuses on minimizing the pixel-wise error between the target image and the fully rendered approximation by brush strokes, it will cause the rendered image to match the target image mainly in color. At the same time MSE loss is prone to blurring which results in washed out edges in the rendered image. Thus, this loss is must be accompanied by additional losses to make up for the shortcoming of MSE loss.

A popular choice for preserving content in an image which, which is associated with preserving edges, is **perceptual loss**. Perceptual loss is based on a VGG Network [VGG] that is pre-trained on ImageNet [ImageNet]. To compute the loss, the activations of deep layers (usually the fourth convolutional block) of the pre-trained VGG network are inferred and then compared using MSE loss.

$$L_{\text{perceptual}} = \frac{1}{H_f, W_f} \sum_{p \in \text{pixels}} \|f(p) - f'(p)\|_2^2 \quad (6.23)$$

The resulting distance is meant to capture how well edges between the two input images are preserved which should be equal to whether the content in both images is the same. Together with MSE loss perceptual loss is often used to get better reconstructions than with MSE loss alone, as edges of object in the image are better preserved, prohibiting blurriness that would occur otherwise.

As an evolutionary step to perceptual loss, Zhang *et al.* [lips] introduced perceptual similarity or **LPIPS loss** which weighs the different layers of the pre-trained VGG-network differently in order to increase the effectiveness of perceptual distance between two images. LPIPS loss is meant to preserve edges even better than perceptual loss does with a similar computational overhead.

Besides losses that operate in pixel space, it is also necessary to restrict the action space for each brush stroke. As it has been explained in section ?? the renderer has been trained on a limited data set which put constraints on how curved brush strokes may be and how the ration between length and width ought to look like. These constraints must be enforced in the optimizing process as well. Because the renderer has nit been trained on data that lies outside of the generated data set, it is very likely that that the renderer will break if the input parameters lie to far apart from the space it has been trained in. The results would then be renderings with

no output, distorted brush strokes or just noisy output, as it can be seen in ??.

Thus, one must think of an additional loss to confine the parameter space to the same space as the generated brush stroke data during optimization. There are two ways of achieving this:

- ▶ Discriminators
- ▶ Explicitly coded losses

Discriminators are a popular choice in this context because even if the data distribution is not known beforehand, a discriminator is still able to learn the distribution from data and thus point out wrong parameter combinations in this case. Still, a discriminator comes with a few compromises, as the target distribution will never be perfectly learned rather than well approximated by a discriminator. This leaves room for weaknesses as well as local minima in the discriminator's prediction which would result in worse quality for this task. Usually, these weaknesses are made up for during adversarial training as if the generator over fits to such weaknesses the discriminator will quickly penalize such a solution. In the case of the optimization routine which is employed for this problem, it is not possible to train the discriminator online as the limit amount of brush strokes will allow the discriminator to over-fit the problem easily. Thus only a pre-trained discriminator with its said weaknesses can be used in this case.

Handpicked losses As the data distribution for the generated data set is actually known in this case (see 6.2), it is also possible to manually define losses that confine the brush strokes. The width constraint – as a first example – can easily be enforced by penalizing whenever the brush stroke's width $w(x)$ is more than half the length $l(x)$ between the start point $\underline{s}(x)$ and the end point $\underline{e}(x)$ of the brush stroke x :

$$L_{\text{bs}} = \frac{1}{|X|} \sum_{x \in X} \max(0, 2w(x) - l(x)) \quad (6.24)$$

$$= \frac{1}{|X|} \sum_{x \in X} \max(0, 2w(x) - \|\underline{s}(x) - \underline{e}(x)\|_2) \quad (6.25)$$

This is a bit more complicated regarding the limitation that is introduced to the control point $\underline{c}(x)$. As $\underline{c}(x)$ was sampled from a multivariate gaussian with fixed parameters in the data set, it should now follow a similar distribution in relation to the direction $\underline{s}(x) - \underline{e}(x)$ of each stroke.

This can be achieved by first defining two orthonormal basis vectors which are either parallel $\underline{n}_{se}^{\parallel}(x)$ or orthogonal $\underline{n}_{se}^{\perp}(x)$ to the directional

vector $\underline{s}(x) - \underline{e}(x)$:

$$\underline{n}_{se}^{\parallel}(x) = \frac{\underline{s}(x) - \underline{e}(x)}{\|\underline{s}(x) - \underline{e}(x)\|_2} \quad (6.26)$$

$$\underline{n}_{se}^{\perp}(x) = R_{\pi/2} \frac{\underline{s}(x) - \underline{e}(x)}{\|\underline{s}(x) - \underline{e}(x)\|_2} \quad (6.27)$$

$$\text{with } R_{\pi/2} = \begin{pmatrix} \cos \pi/2 & -\sin \pi/2 \\ \sin \pi/2 & \cos \pi/2 \end{pmatrix} \quad (6.28)$$

Then $\underline{c}(x)$ can be projected into the coordinate system spanned by $\underline{n}_{se}^{\parallel}(x)$ and $\underline{n}_{se}^{\perp}(x)$:

$$c^{\parallel}(x) = (\underline{c}(x) - \underline{a}(x)) \cdot \underline{n}_{se}^{\parallel}(x) \quad (6.29)$$

$$c^{\perp}(x) = (\underline{c}(x) - \underline{a}(x)) \cdot \underline{n}_{se}^{\perp}(x) \quad (6.30)$$

$$\underline{a}(x) = \frac{\underline{s}(x) + \underline{e}(x)}{2} \quad (6.31)$$

Now, the axes of the original multivariate distribution co-align with $c^{\parallel}(x)$ and $c^{\perp}(x)$. By calculating the mean and standard deviation along these projections, they can be compared to the parameters of the original data distribution.

$$\underline{\mu} = \frac{1}{|X|} \sum_{x \in X} \begin{pmatrix} c^{\parallel}(x) \\ c^{\perp}(x) \end{pmatrix} \quad (6.32)$$

$$\underline{\Sigma} = \left(\frac{1}{|X|} \sum_{x \in X} \begin{pmatrix} c^{\parallel}(x) \\ c^{\perp}(x) \end{pmatrix} - \underline{\mu} \right) \left(\begin{pmatrix} c^{\parallel}(x) \\ c^{\perp}(x) \end{pmatrix} - \underline{\mu} \right)^T \frac{1}{2} \quad (6.33)$$

$$(6.34)$$

Using the Kullback-Leibler divergence for multivariate normal distributions, the compliance with the data sets distribution can be checked:

$$\mathcal{L}_{KL} = \frac{1}{2} \left[\log \frac{|\Sigma|}{|\tilde{\Sigma}|} - d + \text{tr}(\Sigma^{-1} \tilde{\Sigma}) + (\mu - \tilde{\mu})^T \Sigma^{-1} (\mu - \tilde{\mu}) \right] \quad (6.35)$$

with

$$\tilde{\mu} = \begin{pmatrix} 0 \\ 0 \end{pmatrix} \quad (6.36)$$

$$\tilde{\Sigma} = \begin{pmatrix} \frac{1}{200} & 0 \\ 0 & \frac{1}{25} \end{pmatrix} \text{ see (6.4)} \quad (6.37)$$

since the origin of the projection is at $\underline{a}(x)$, thus in the middle between $\underline{s}(x)$ and $\underline{e}(x)$ and coincides with the center of the data distribution.

In theory, L_{KL} and L_{bs} should be able to capture any deviation from the source data and ensure that parameters stay within the training space of the renderer. One problem that obviously could arise with this formulation is for patches with very similar brush strokes that show a mean other than $\tilde{\mu}$ as well as a very low values inside the covariance

matrix and non-zero off-diagonal values. Thus it will be favorable to include as many strokes as possible when calculating this loss, ideally all strokes in the global parameter catalogue.

Style Transfer

One question that arises when talking about the optimization procedure, is whether there could be style transfer performed with this approach as well. Especially, since the optimization procedure has explicitly been compared to the approach by Gatys et al. [4], it is just a short way of thought to assume that such an approach could also be applicable in this case.

This mainly requires to introduce a style loss as it has been done by Gatys et al. [4], since a content loss is already in place (see (6.23)). A style loss can then be implemented similarly by aggregating the activations of more layers and then calculating the gram matrices:

$$\mathcal{L}_{\text{style}} = \sum_{l=0}^L \frac{w_l}{4N_l^2 M_l^2} \sum_{ij} (G_{ij}^l - A_{ij}^l)^2 G_{ij}^l = \sum_k F_{ik}^l F_{jk}^l \quad (6.38)$$

Optimization Details

As it has been explained in Section ??, it is not possible to optimize brush strokes for the whole image in parallel. This is partly due to the memory requirements of the neural renderer which scales with the number of rendered patches. And partly due to the memory requirements of the placing and blending method, which scales with the number of rendered brush strokes as well as the number of surrounding fixed brush strokes and the patch window size. Experiments have shown that a combination of 256 rendered brush strokes with 128 surrounding brush strokes paired with a patch window size of 320x320 pixels occupies around 9.5GiB of memory which leaves enough space for more advanced losses and tweaks to the network architecture.

The learning rate for the optimization procedure can be significantly larger than for training the neural renderer. With a learning rate of 0.01 the optimization procedure will converge significantly faster to a solution without any instability issues.

One choice, which has to be made individually per target image is how many brush strokes will cover the image. As larger images obviously require more brush strokes than smaller images, what should remain the same is the **brush strokes density**. The brush stroke density will decide how many pixel on average should be covered by each brush stroke and thus be used during initialization. A number of $100 \frac{\text{pixels}}{\text{brush stroke}}$ has produced the best results during the experiments.

Another important choice that goes along the choice of how dense the brush strokes should be distributed, is that of how many optimization steps each brush stroke will be object to. As for too few steps the training will not have converged and for too many optimization will take an unnecessary amount of time. As this can vary between images, since some images require more time to converge, about $1500 \frac{\text{steps}}{\text{brush stroke}}$ give

constantly good results. There is no direct enforcement that each brush strokes is updated exactly this often but it can be expected that due to the uniform sampling of render patches, there will be no major deviation for some brush strokes. Since 256 brush strokes will be optimized in every step, the total number of optimization steps can be calculated together with the brush stroke density and the image's size.

Another minor detail is the fact that for each render patch five consecutive optimization steps are performed as this safes memory bandwidth, since data must be written and read from memory each time a different render patch is optimized.

Results

Results can be seen in Figure ?? for Starry Night as the standard reference image of this thesis. It took approximately 15,000 optimization steps which corresponds to about 2h. The rendering consists of roughly 10,000 brush strokes.

Figure ?? shows the result for natural photo as target image which took also 2h to compute with 8.000 brush strokes and 12.000 optimization steps.

7

Ablation Experiments

In this chapter, the ablation experiments for the approaches presented in Chapter ?? shall be presented. This is meant to show some weaknesses of each method as well as give a visual understanding of the effects of some losses and parameters. Together with an interpretation it should help to get a better understanding of previous motivations.

7.1 Neural Renderer	41
7.2 Optimization Procedure	41

7.1 Neural Renderer

The neural renderer will be evaluated with respect to a few aspects. First, the performance for larger resolution images will be evaluated. Then, the effect of the discriminator will be shown by training the renderer without the discriminator loss. At last, the weakness of the renderer, which is data points outside the training space, will be analyzed.

7.2 Optimization Procedure

The optimization procedure gives many starting points for ablation experiments. At first, the effects of each loss will be visually compared. Then, different initializations will be shown as well as various choices for training parameters.

CONCLUSION

Conclusion

8

Outlook

9

APPENDIX

Appendix A

Bibliography

Here are the references in citation order.

- [1] Reiichiro Nakano. 'Neural Painters: A learned differentiable constraint for generating brushstroke paintings'. In: (2019) (cited on page 17).
- [2] Zhewei Huang, Wen Heng, and Shuchang Zhou. 'Learning to Paint With Model-based Deep Reinforcement Learning'. In: (2019) (cited on page 17).
- [3] Yaroslav Ganin et al. 'Synthesizing Programs for Images using Reinforced Adversarial Learning'. In: (2018) (cited on page 18).
- [4] Leon A Gatys, Alexander S Ecker, and Matthias Bethge. 'A Neural Algorithm of Artistic Style'. In: (2015) (cited on page 39).

List of Figures

6.6 Visualization of the generator and discriminator architectures	23
--	----

List of Tables